

CSCI 335

Software Design and Analysis

III

Splay Trees/B-trees

Chapter 4

Ioannis Stamos

Amortized cost

- Consider a sequence of M operations(insert/delete/find)
 - What is the total cost?
 - What is the average cost per operation?

Amortized cost

- Consider a sequence of M operations (insert/delete/find):
- Example: binary search tree (regular)
 - M operations can cost in the worst case $O(M * N)$
 - Each operation will **not** cost more than $O(N)$
 - On average each operation costs $O(N)$

Amortized cost

- Consider a sequence of M operations (insert/delete/find):
- AVL tree:
 - A sequence of M operations will cost $O(M * \log N)$
 - Each operation will **not** cost more than $O(\log N)$
 - On average each operation costs $O(\log N)$

Amortized cost

- Consider a sequence of M operations (insert/delete/find):
- Suppose that total cost is $O(M * f(N))$ irrespective of the actual sequence of operations.
- The average cost is $O(f(N))$ for each operation.
- This is called **amortized running time**.
- Caveat:

Individual operations in the sequence can be expensive though !

Splay tree

- A tree with amortized running time of $O(\log N)$
- Some operations could be more expensive
- How can we achieve this?

Splay tree

- The trick is to rebalance the tree after a **find()** operation
- Bring the item returned by find() to the root while applying rotations on the way to the root.


Splay tree

- Successive **finds()** of same element will be pretty fast
- Other items are coming **closer** to the root
- No need to store height information at each node
- **Amortized cost** of sequence of M operations is $O(\log N)$


Splay Trees: a simple idea that does not work


- After `find()` perform single-rotations bottom-up

Single rotation (case 1)




Single rotation (case 4)








The access path is dashed. First, we would perform a single rotation between k_1 and parent, obtaining the following tree.




Then, we rotate between k_1 and k_3 , obtaining the next tree.




Then two more rotations are performed until we reach the root.






then two more rotations are performed until we reach the root.



K_1 went up



K_3 went down

Splay Trees:

a simple idea that does not work

- There is a sequence of M operations requiring $\Omega(N)$ time (**amortized**).
- \Rightarrow we want logarithmic amortized time

Splay Trees:

M operations requiring $\Omega(N)$ time




- Consider inserting 1,2,3, ... , N into an initially empty tree
 - Note that you splay on insertion (i.e. single AVL rotation)=>
 - Only left children
- Total time to build tree is $O(N)$ (not bad)

Insert 1, 2, 3, 4, 5, 6, ...



Insert 1






Insert 2




Insert 3




Insert 4




...



Insert 1, 2, 3, 4, 5, 6, ... , N




Total time to build tree is $O(N)$




Access 1, 2, 3, 4, 5, 6

Access 1




Access 1, 2, 3, 4, 5, 6

Access 2



Access 3



Splay Trees:

M operations requiring $\Omega(N)$ time

- Accessing the sequence 1,2,... is $\Omega(N^2)$ though...
- Run example:
 - Access 1 (time $O(N)$), then perform sequence of single rotations (time $O(N)$)
 - Access 2 time $O(N-1)$...

Splaying – the correct way!

Rotate bottom-up on access, along access path

Let **X** a non-root node on access path at which we are rotating
if (parent of **X** is the root) { single rotation with root }

else {

Let **P** be parent of **X**, and **G** be the grandparent

Two cases (plus symmetries) to consider:


- Case 1 (zig-zag): double rotation (see figure)
- Case2 (zig-zig): (see figure)

}


Applet:


<https://www.cs.usfca.edu/~galles/visualization/SplayTree.html>

Zig-zag




Zig-Zig







The first splay step is at k_1 , and is clearly a zig-zag, so we perform a standard AVL double rotation using k_1 , k_2 , and k_3 . The resulting tree follows.




The next splay step at k_1 is a zig-zag, so we perform a standard AVL double rotation using k_1 , k_2 , and k_3 .




The next splay step at k_1 is a zig-zig, so we do the zig-zig rotation with k_1 , k_4 , and k_5 , obtaining the final tree.



Algorithm 16.10: Splay operation



The next splay step at k_1 is a zig-zig, so we do the zig-zig rotation with k_1 , k_4 , and k_5 , obtaining the final tree.




Roughly halving depth
of most nodes...


Algorithm 16.10: Splay (continued)

Example


Insert 1,...,7 -> then access 1



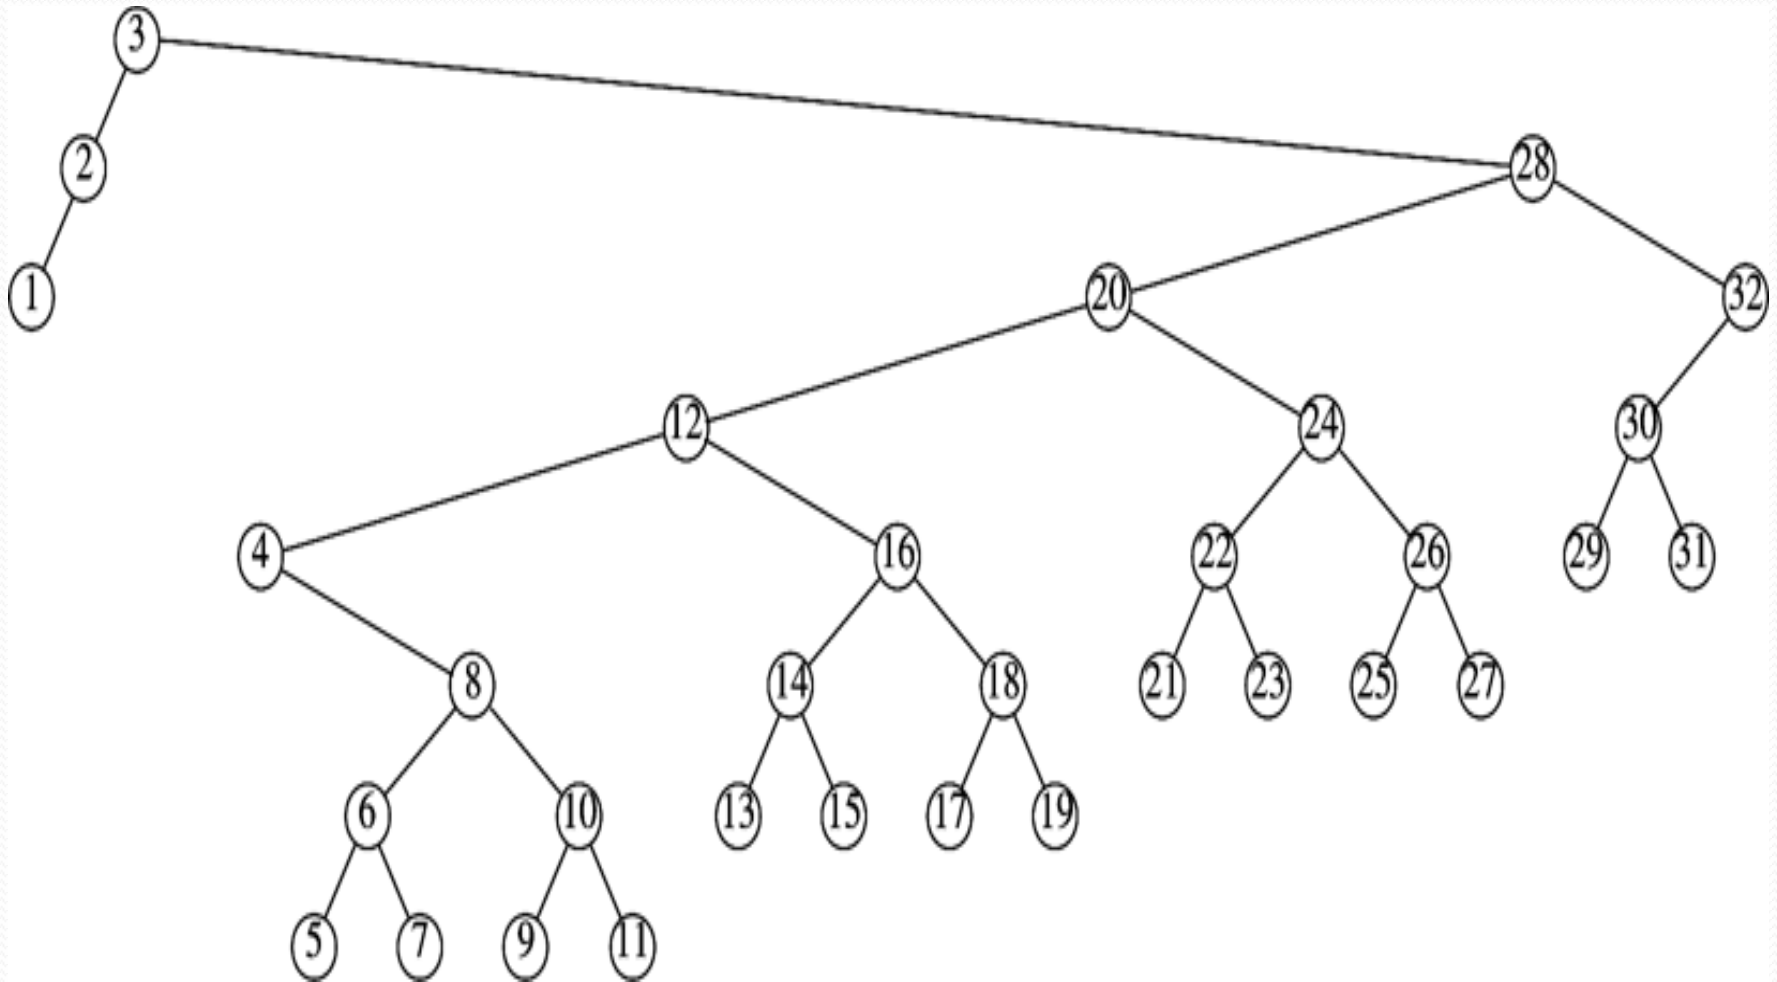
Same example but w/ 32 nodes




Same example but w/ 32 nodes




Same example but w/ 32 nodes




Same example but w/ 32 nodes




Same example but w/ 32 nodes




Same example but w/ 32 nodes




Same example but w/ 32 nodes



Same example but w/ 32 nodes



Same example but w/ 32 nodes



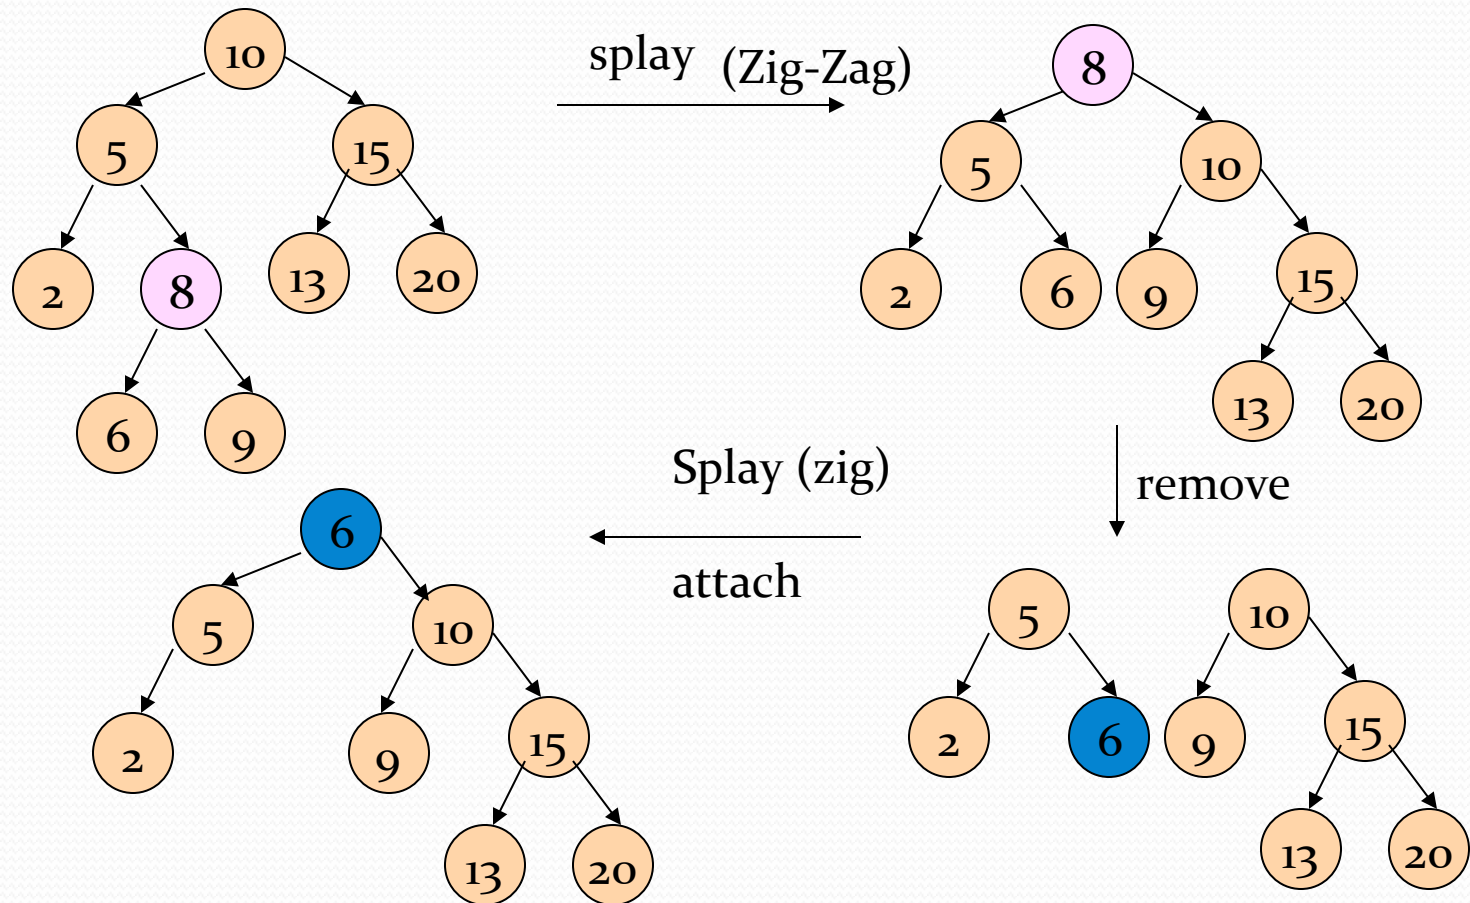
Analysis

- Hard (see later chapter)
- In the example: Access of 2 \rightarrow $N/4$ of root, ..., up-to $\log N$ of root
- Fundamental properties:
 - When access paths are long, longer search time, but rotations good for future operations
 - When access is cheap, rotations not as good.
 - It can be proved that time is $O(\log N)$ per operation (amortized)
- Deletion?
- Much simpler to program with fewer cases
- No need to store balance information

Deletion

- Access, bring node to top
- Delete creating two subtrees
- Access largest element in left tree, bring it to root with no right child
- Patch in right tree as right child

Example Deletion of 8




Exercise from book

- 4.26 (Double rotation implementation without two single rotations)
- 4.27, 4.28

For exercises


Find: 3, 9, 1, 5
Delete: 6



B-trees

- Scenario: Tree is large and can't fit into memory
- Fact: Disk I/O is much slower than machine instruction (one disk access $\sim 4M$ instructions)
- Assume that we have 10M records, each of 256 bytes, and key is 32 bytes
- Solution?
 - AVL? (worst case is $1.44 \log N$, but each operation is costly)
 - We need even smaller trees
 - M-ary tree has height $\log_M(N)$

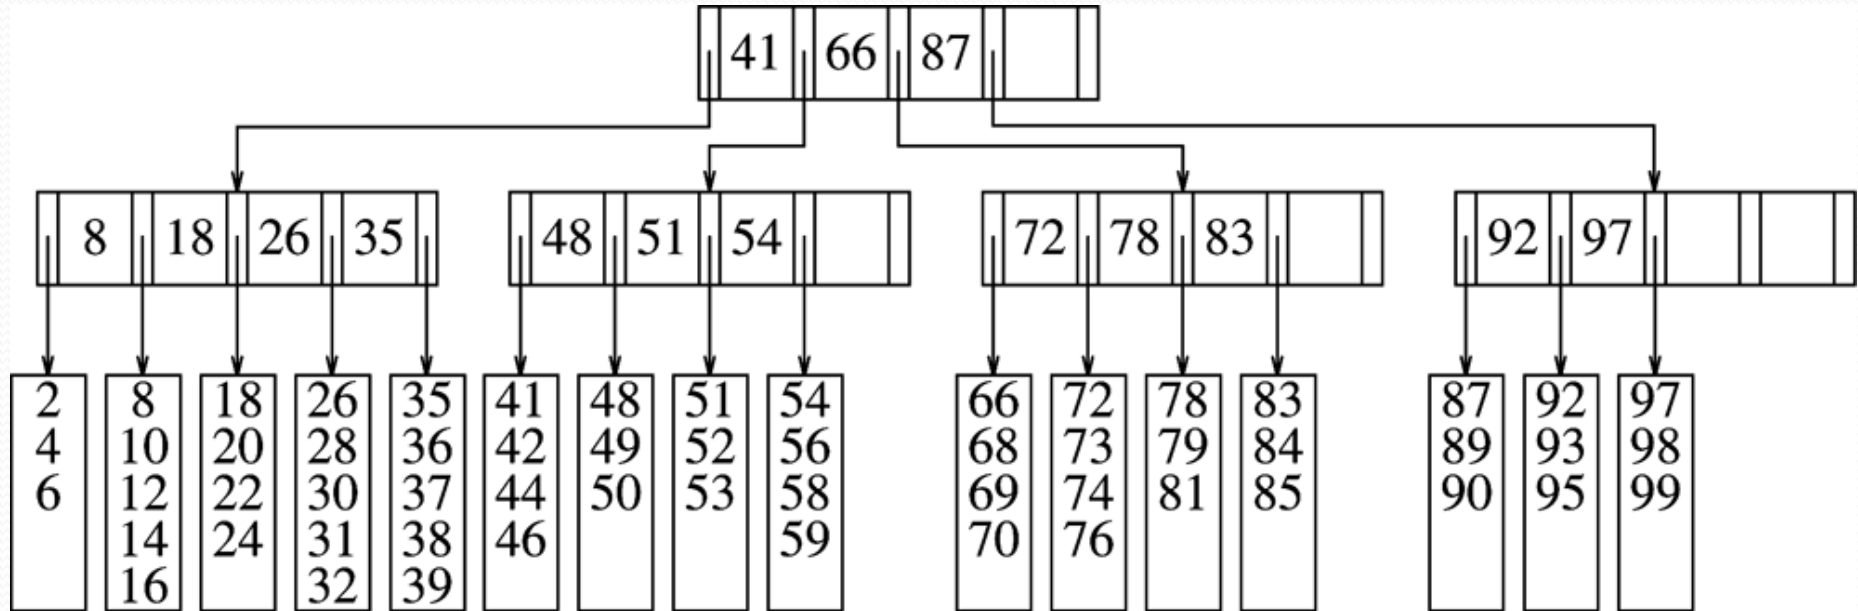
5-ary tree of 31 nodes



Definition of B-tree

- M-ary tree such that:
 - **Data** is stored at leaves
 - The **nonleaf** nodes store up to **$M-1$ keys** to guide searching. Key i represents smallest key in subtree $i+1$
 - **Root** is either a leaf or has between two and M children
 - All nonleaf nodes (except root) have between **$M/2$ and M** children (up to half full).
 - All leaves are at same depth and have between **$L/2$ and L** data items for some L (up to half full).
- **M , and L** are determined based on disk block (one access should load a whole node).

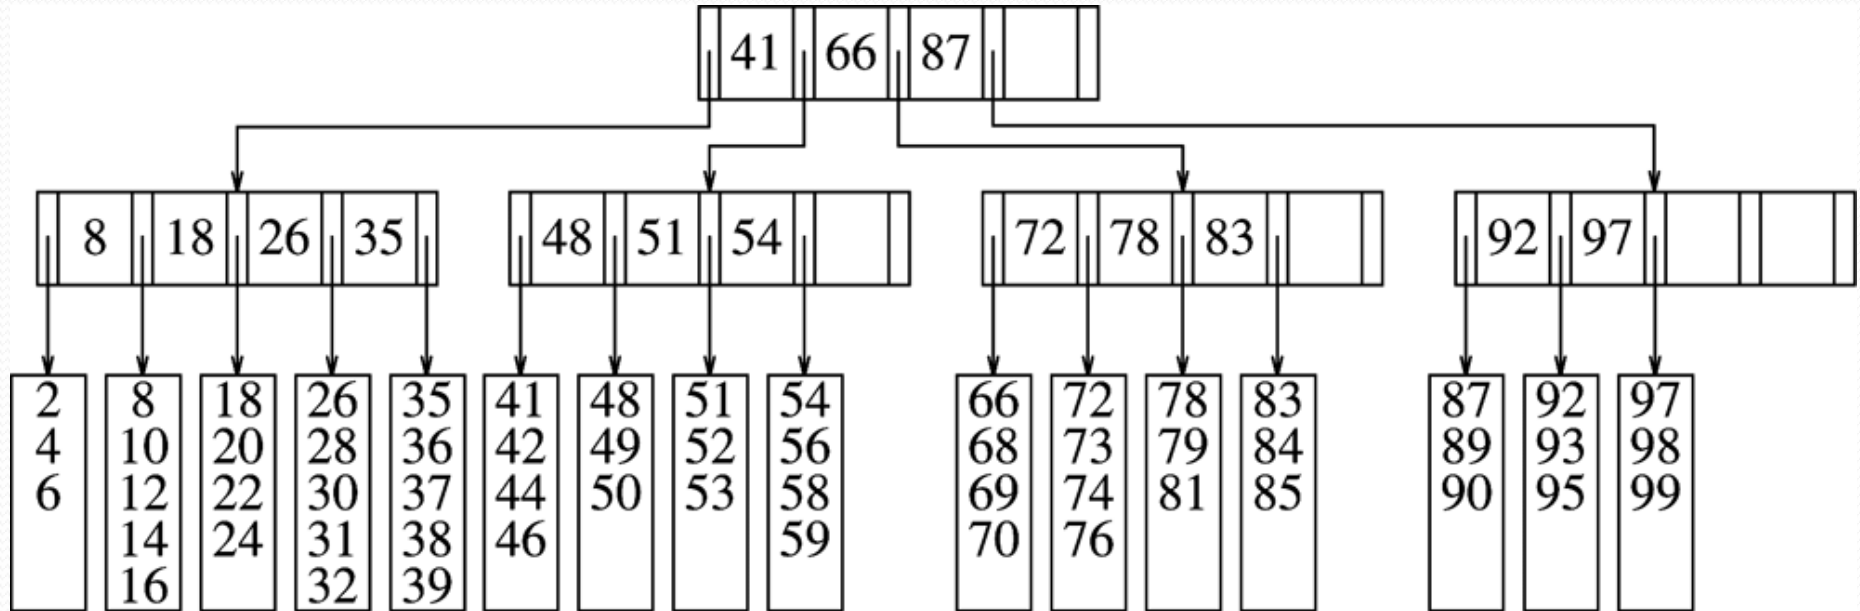
Example B-tree of order 5 ($M=5$)



Example

- **Block** size is 8,192 bytes
- **Key** is 32 bytes
- Internal nodes hold $M-1$ keys $\Rightarrow 32(M-1)$ bytes plus M branches (4 bytes per branch) \Rightarrow total is $36M-32$ bytes.
- Since $8,192 \leq 36M-32 \Rightarrow M = 228$
- Data record is 256 bytes $\Rightarrow 32$ records on a block $\Rightarrow L=32$
- So each leaf should hold between 16 to 32 records
- Each internal node should have **at least 114** branches
- 10 million records $\Rightarrow 625,000$ leafs (10 million/16). In worst case leaves on level 4 (why?)
 - **On average number of accesses is $\log_{M/2}(N)$**
 - Root and first levels could also be cached in memory...

Example B-tree of order 5

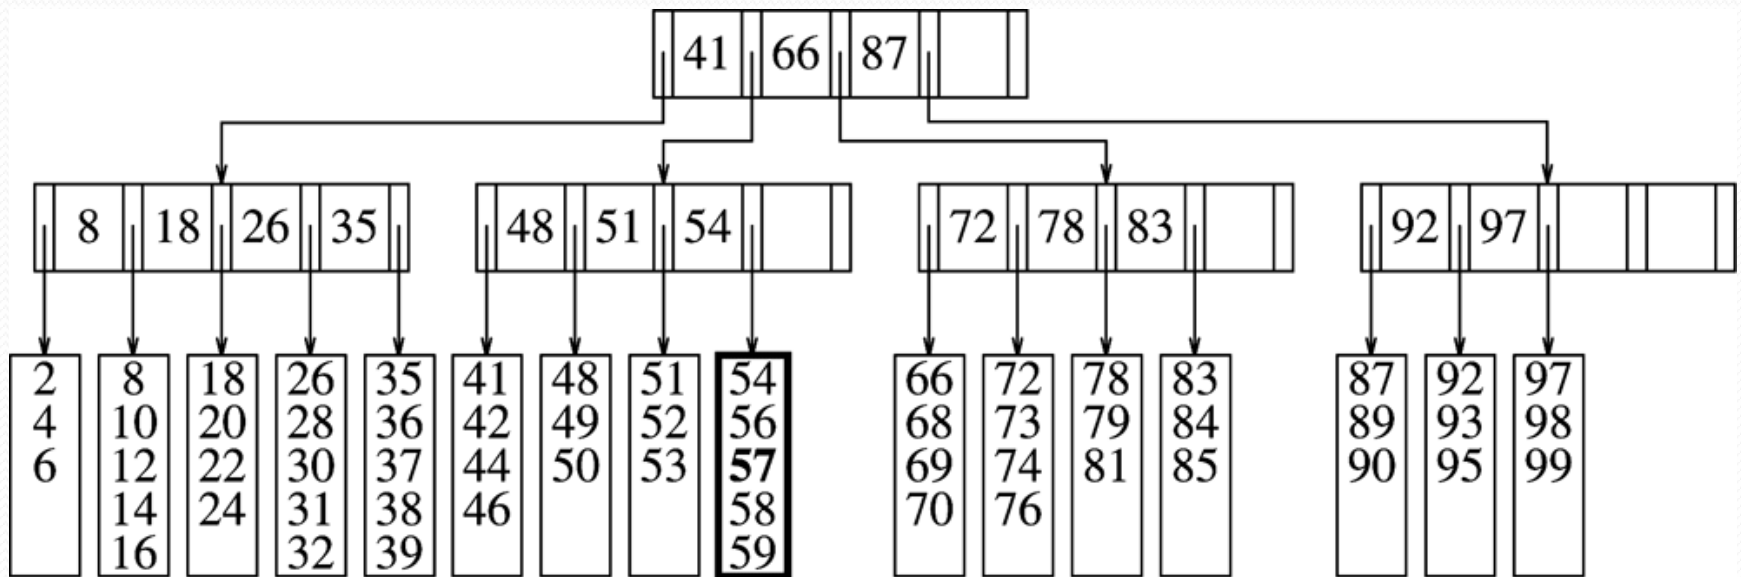


Insertion

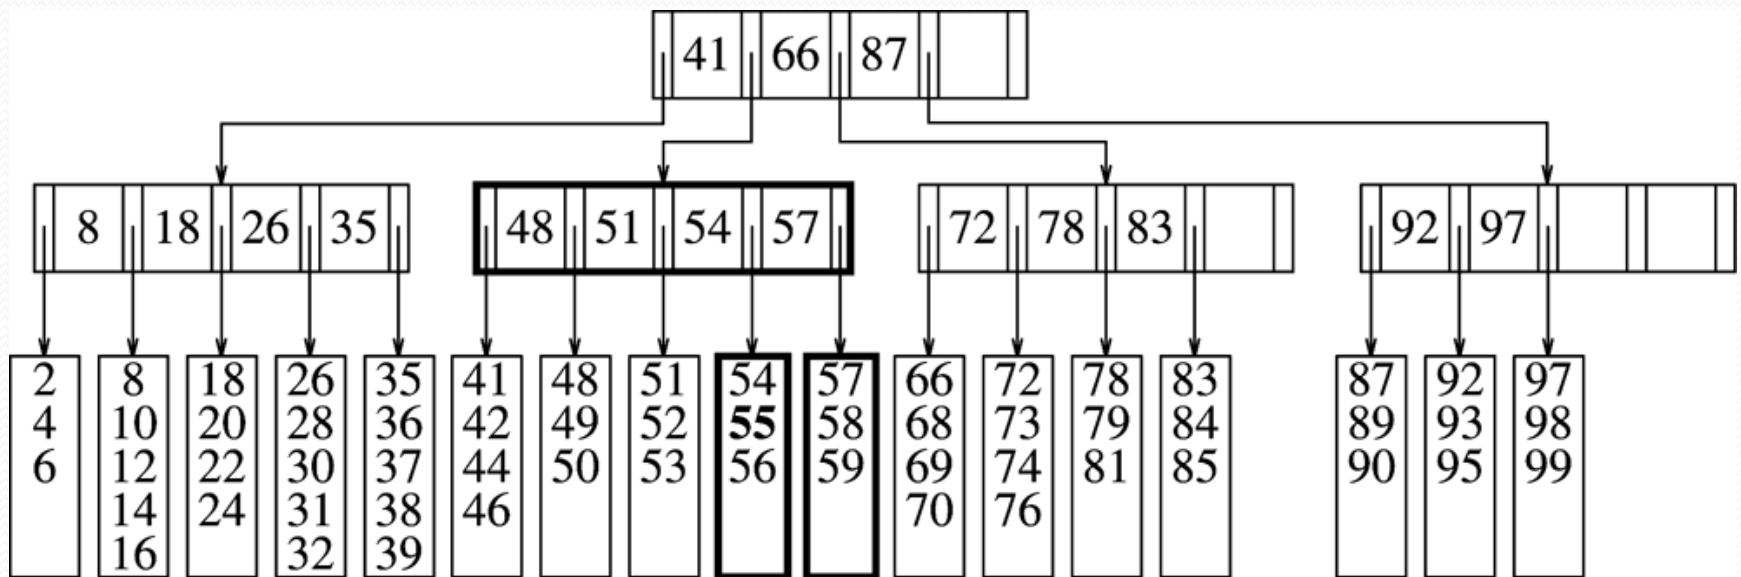
To insert:

- Put it in the appropriate leaf.
 - If the leaf is full, break it in two, adding a child to the parent.
 - If this puts the parent over the limit, split upwards recursively.
 - If you need to split the root, add a new one with two children.
- This is the only way you add depth.

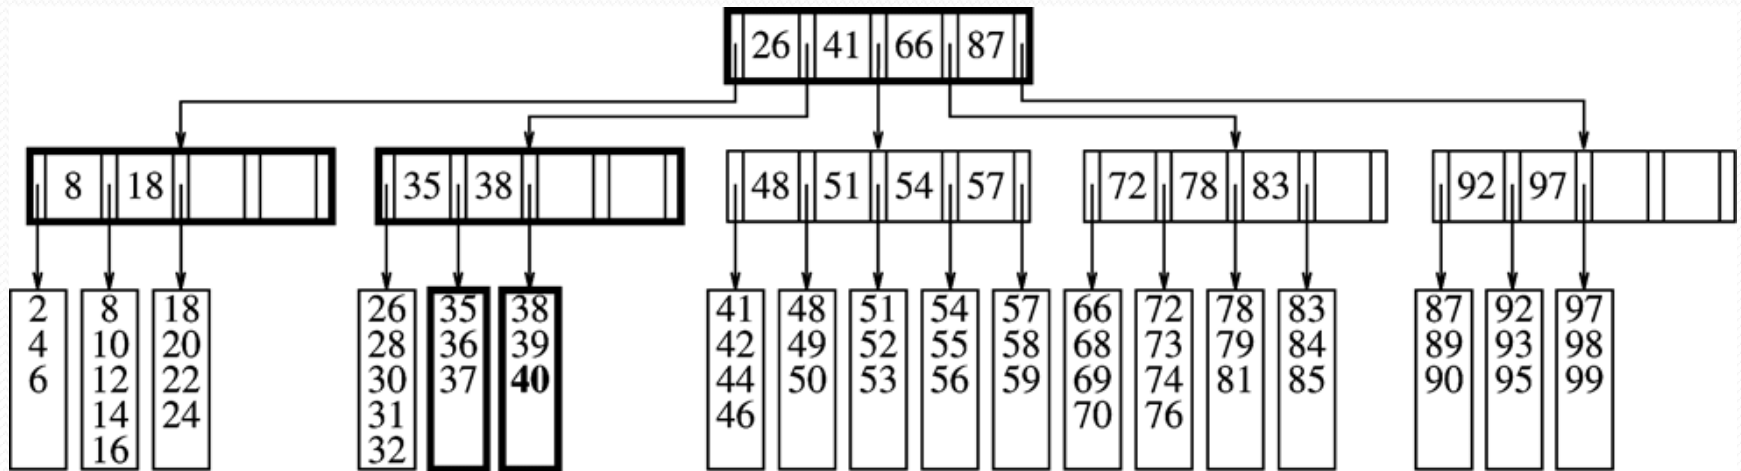
Example of insertion (57)



Example of insertion (55)



Example of insertion (40)



deletion

Delete from appropriate leaf.

- If the leaf is below its minimum, adopt from a neighbor if possible.
- If that's not possible, you can merge with the neighbor. This causes the parent to lose a branch and you continue upward recursively.

Example of deletion (99)

